Lecture 5: Attention, transformer architecture

Leonid Sanochkin

RNN drawbacks

I arrived at the bank after crossing the ... ...street? ...river?

What does **bank** mean in this sentence?

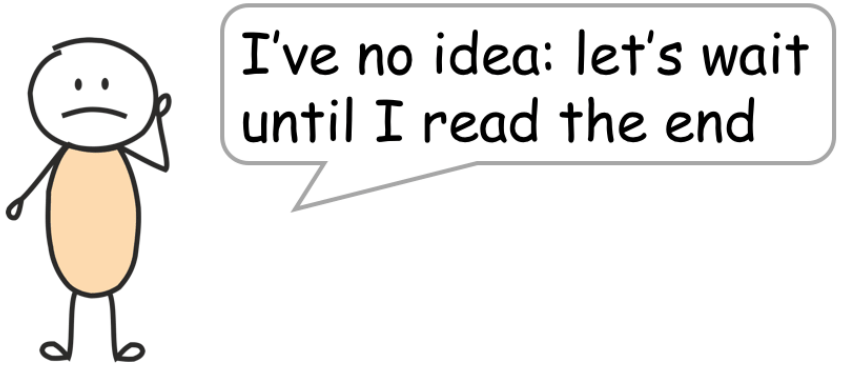

RNNs

O(N) steps to process a sentence with length N

Transformer

Constant number of steps to process any sentence

Self-attention

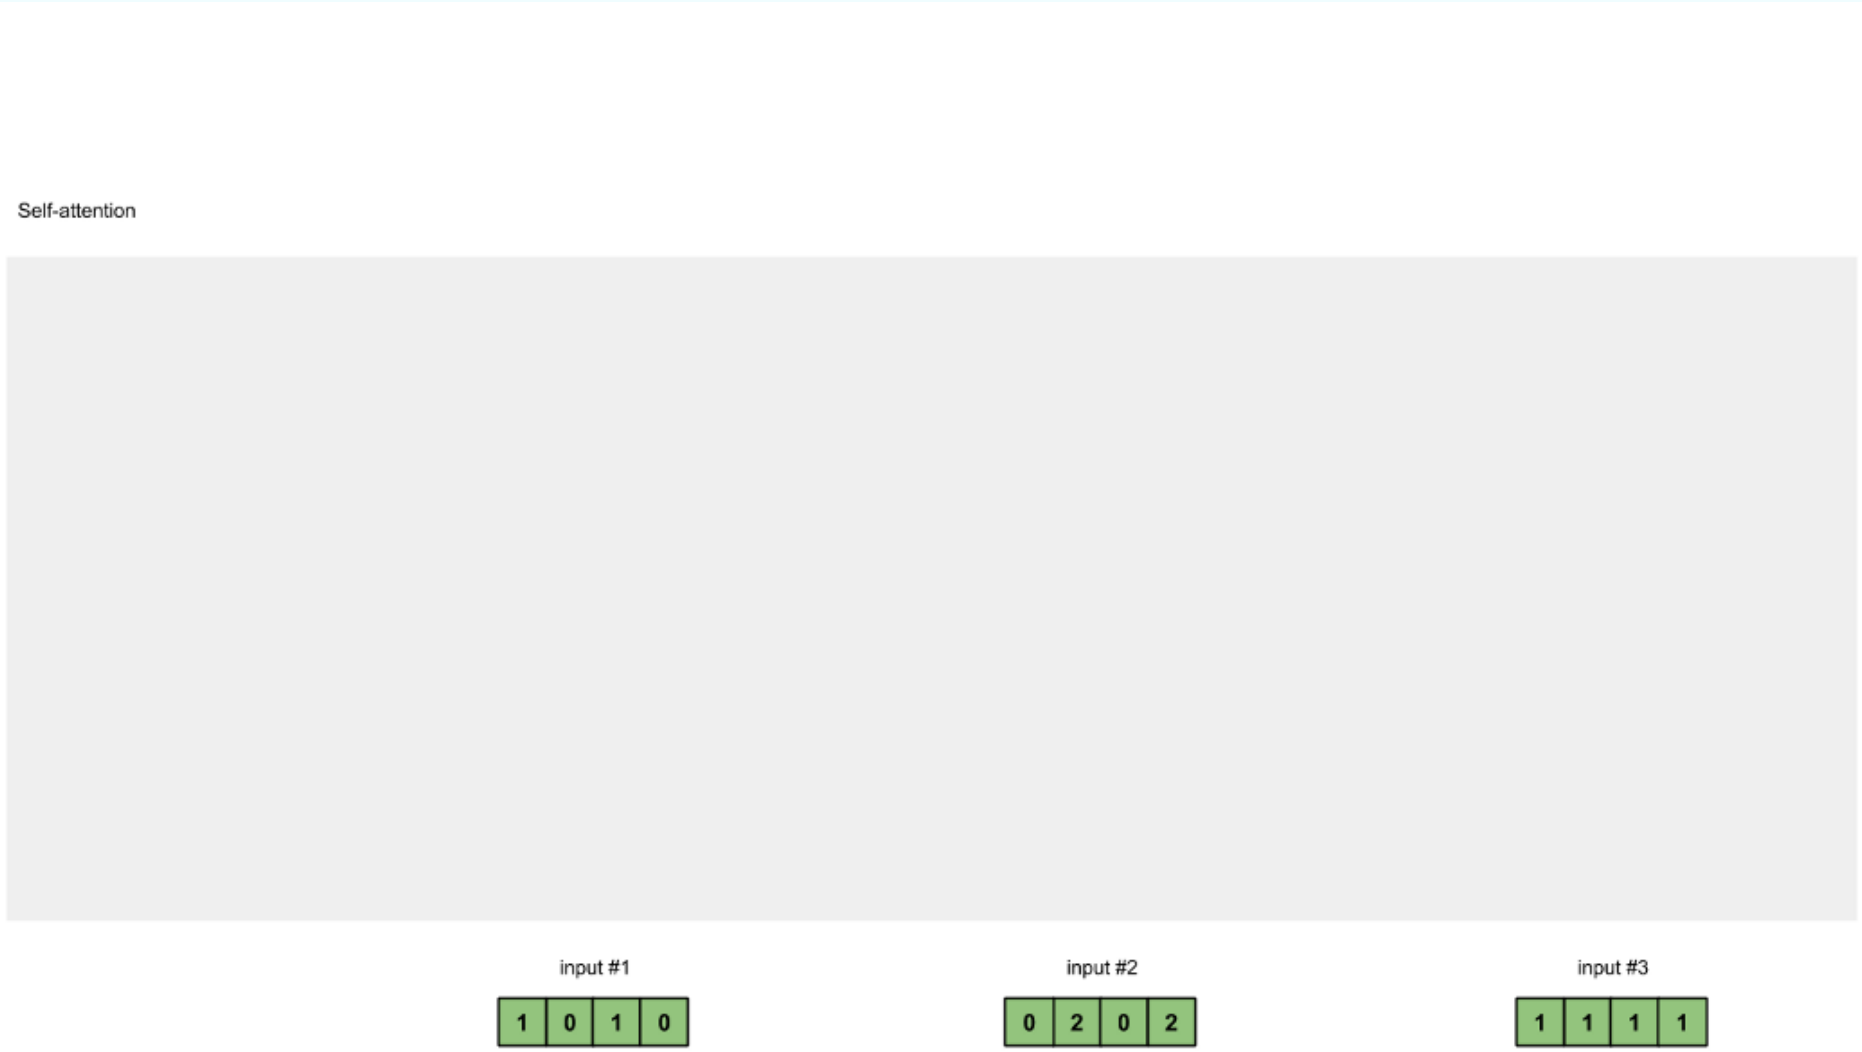

Self-attention

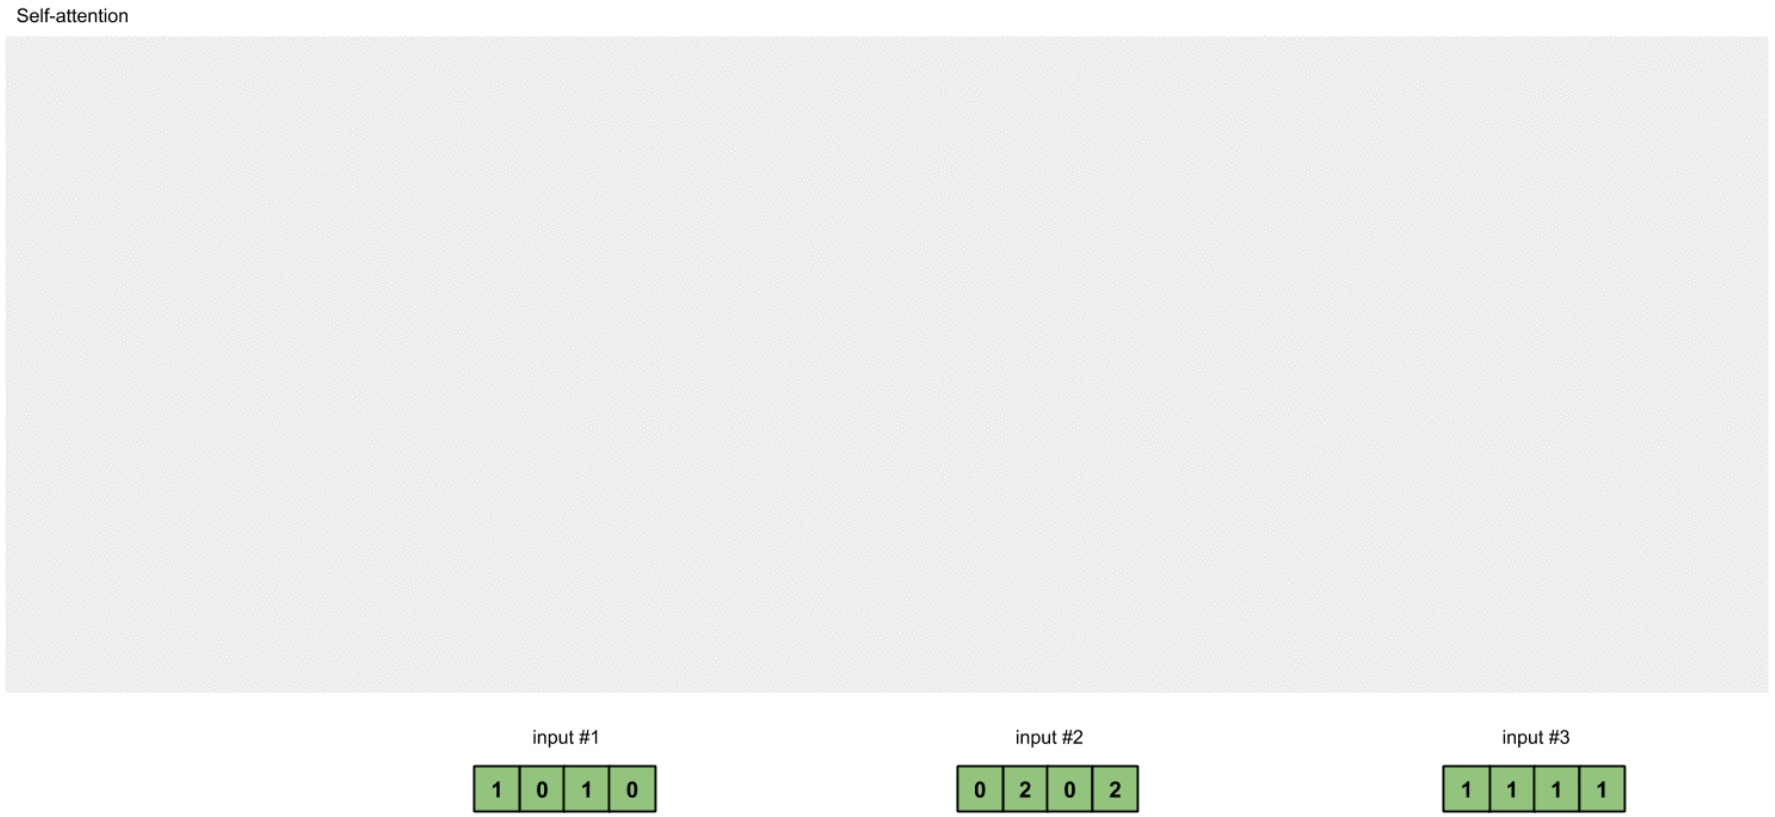

Self-attention

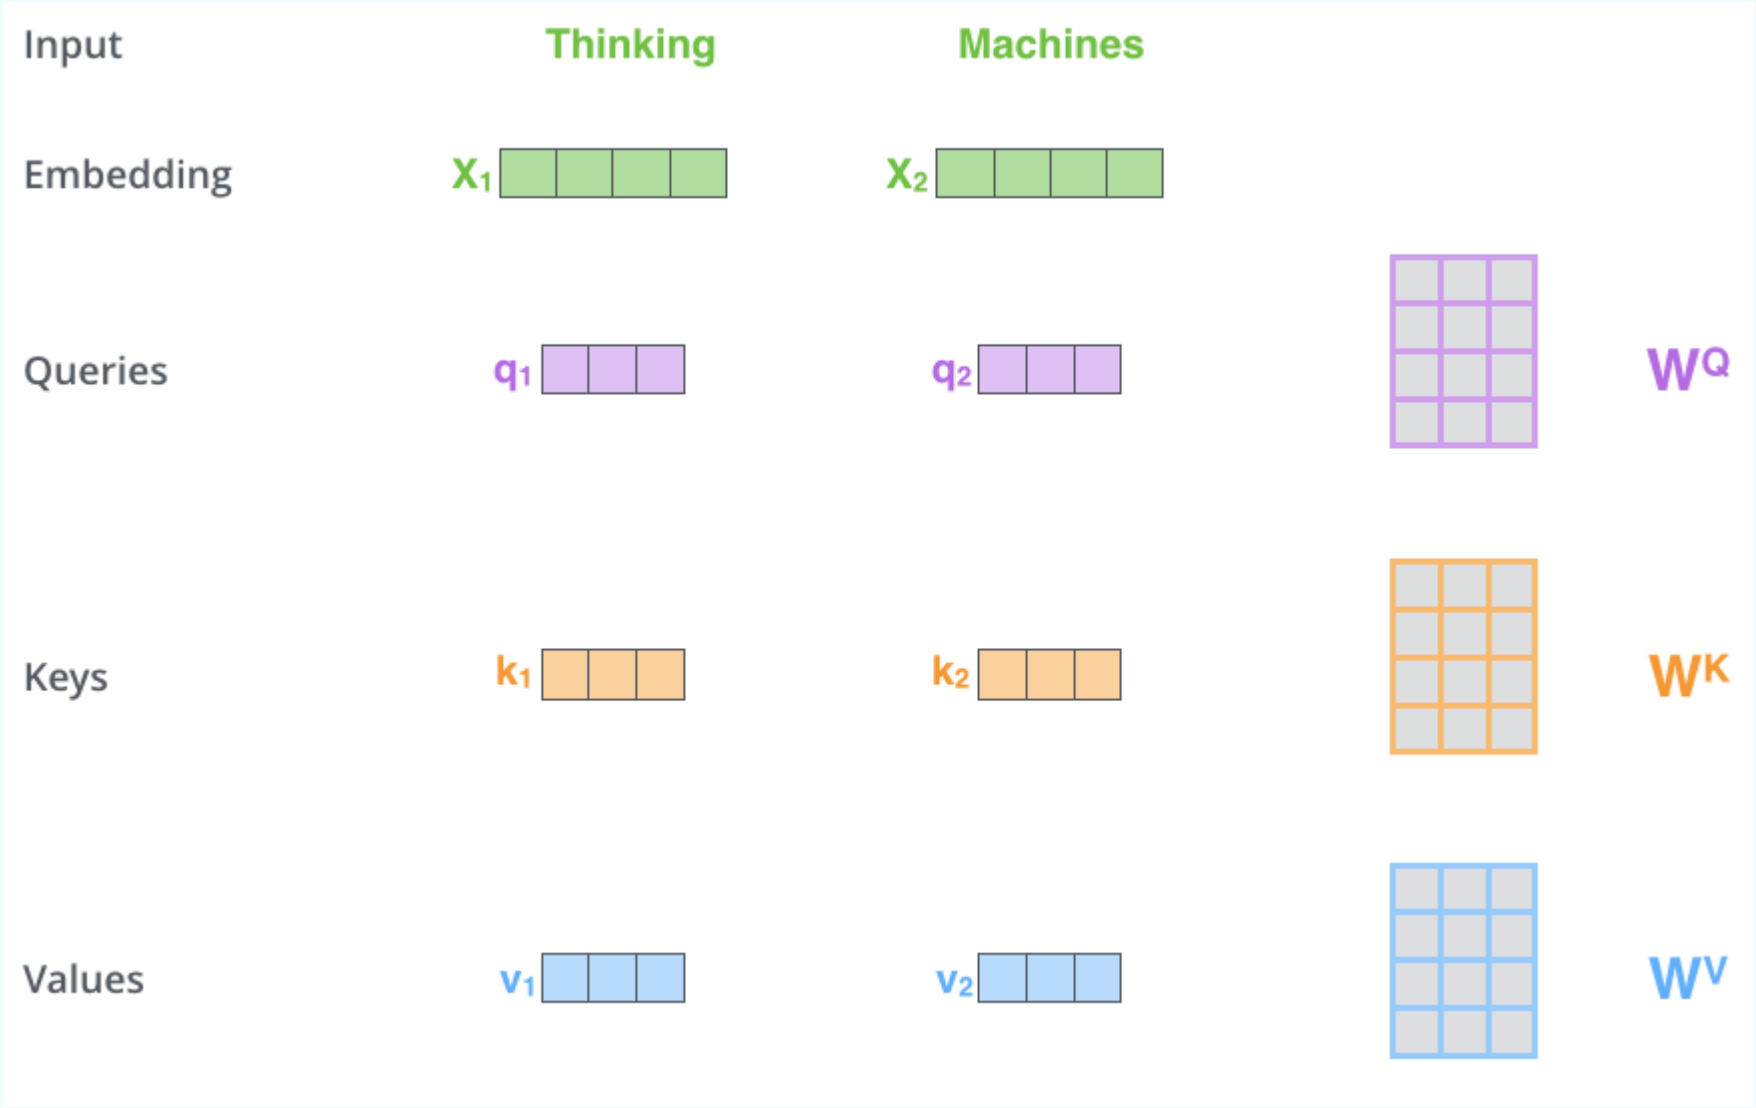

Self-attention

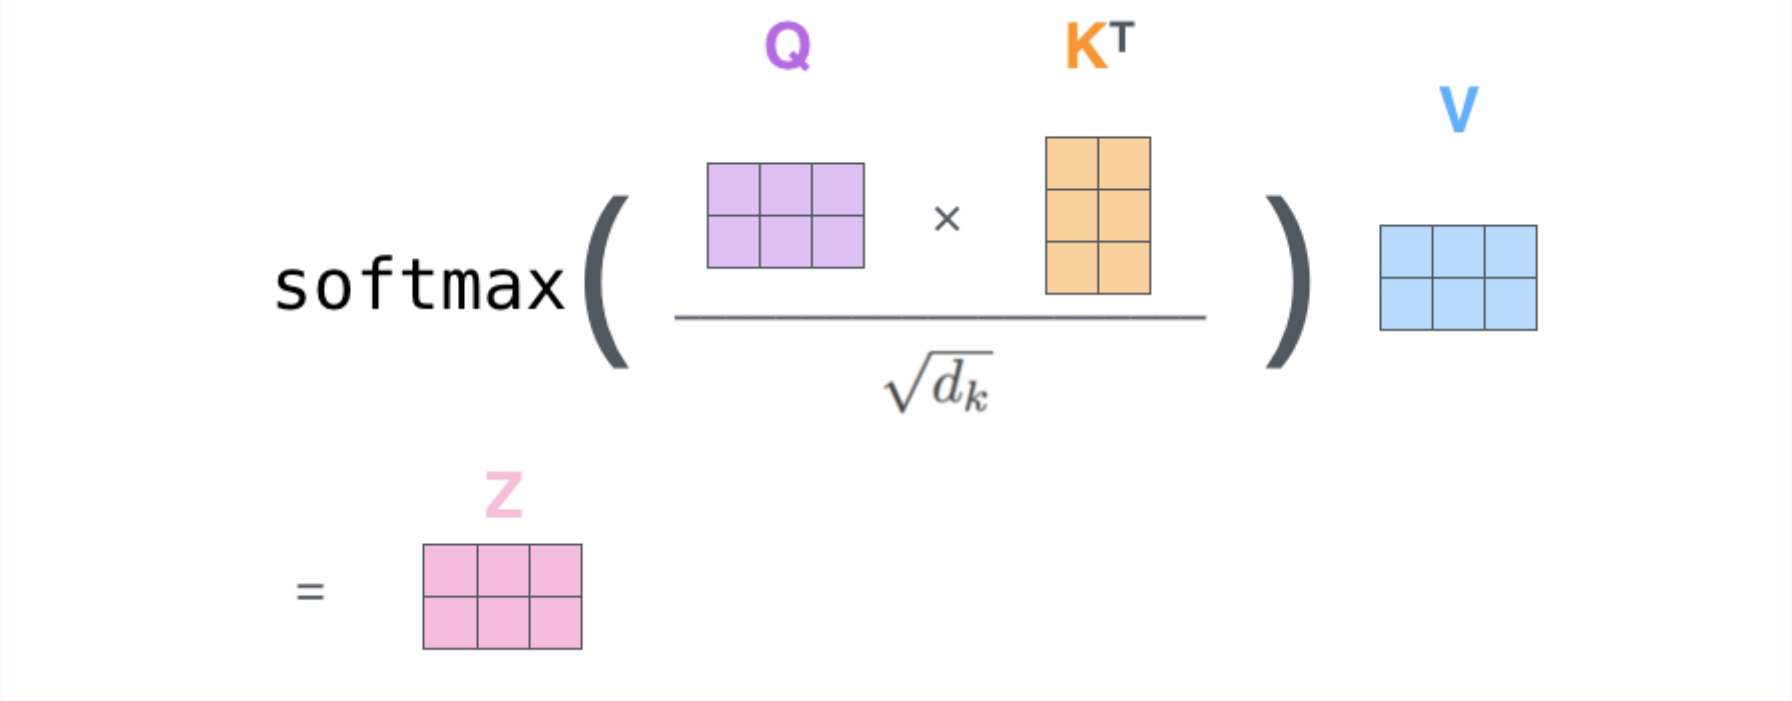

Self-attention

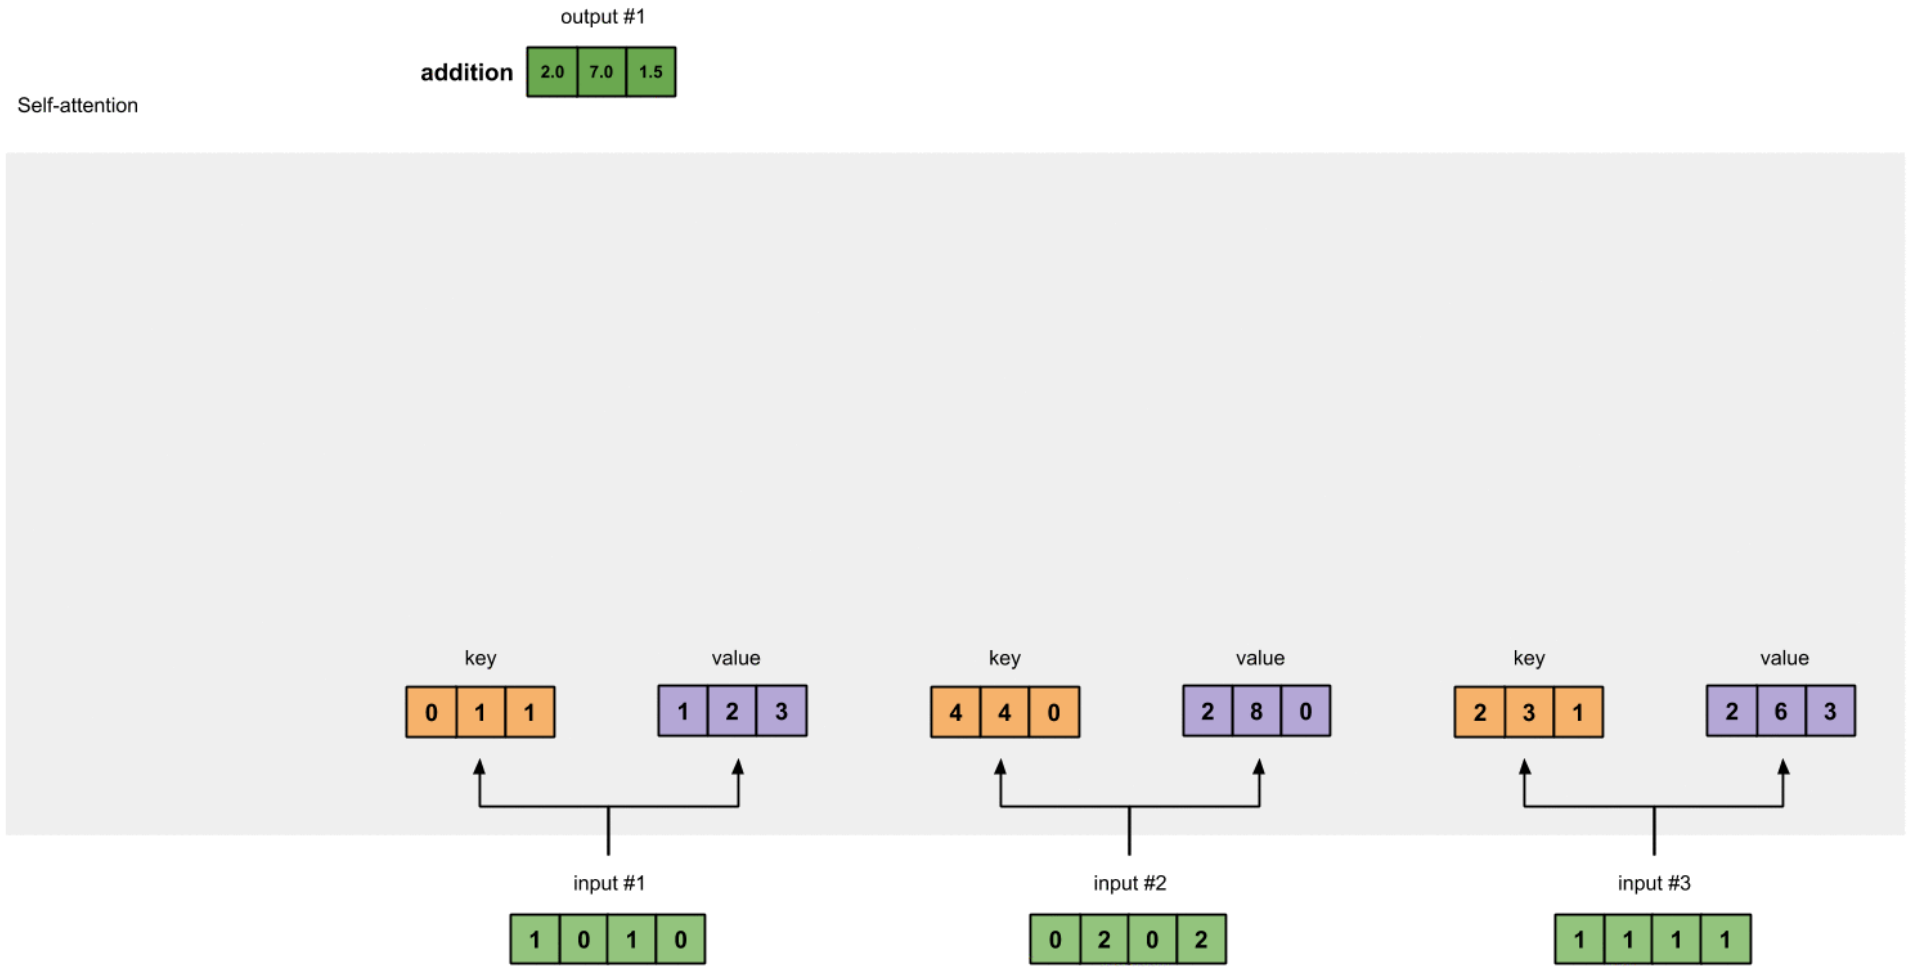

Self-attention

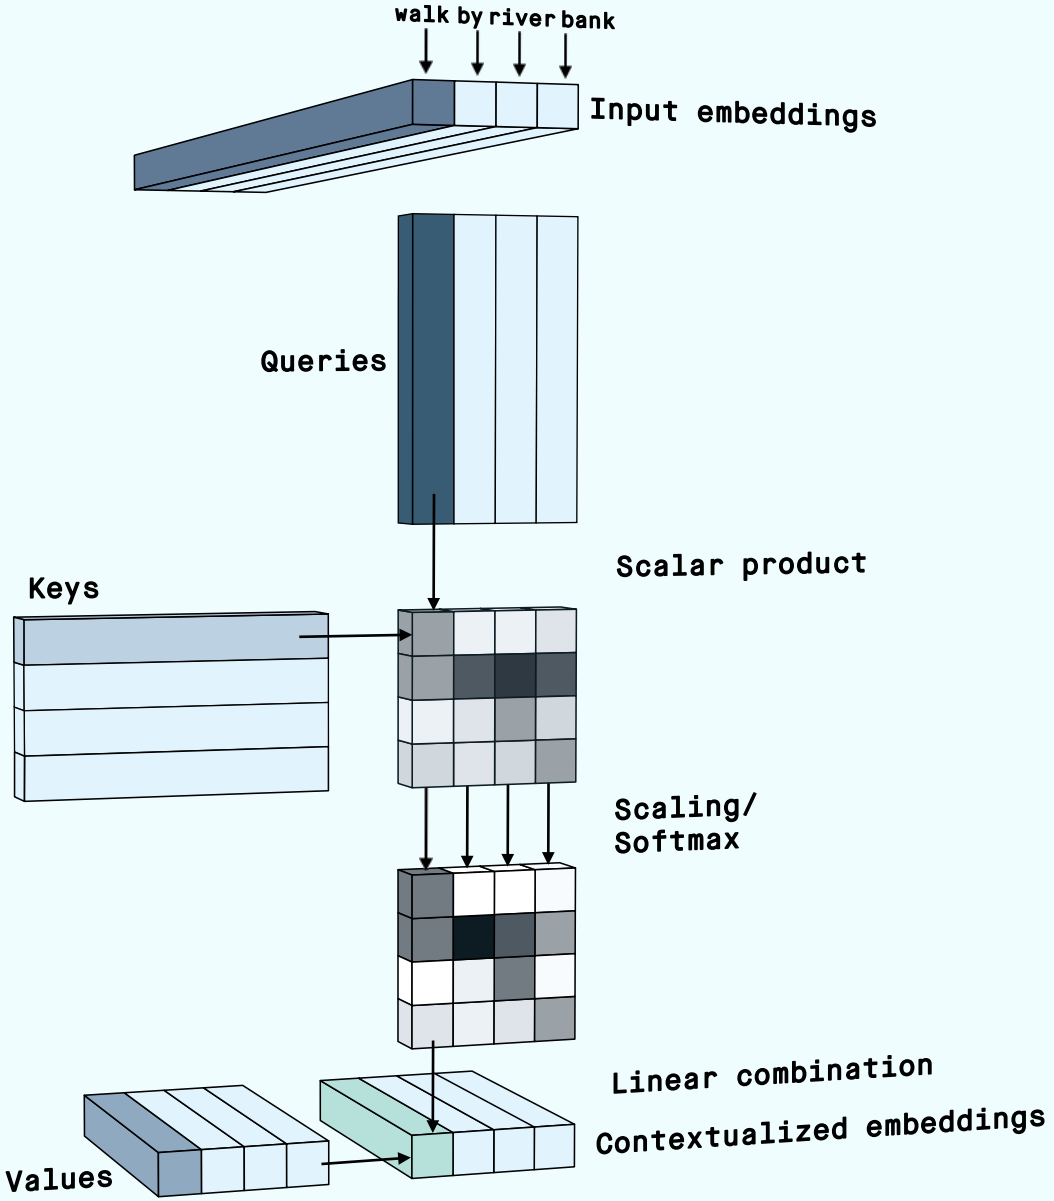

Multihead self-attention

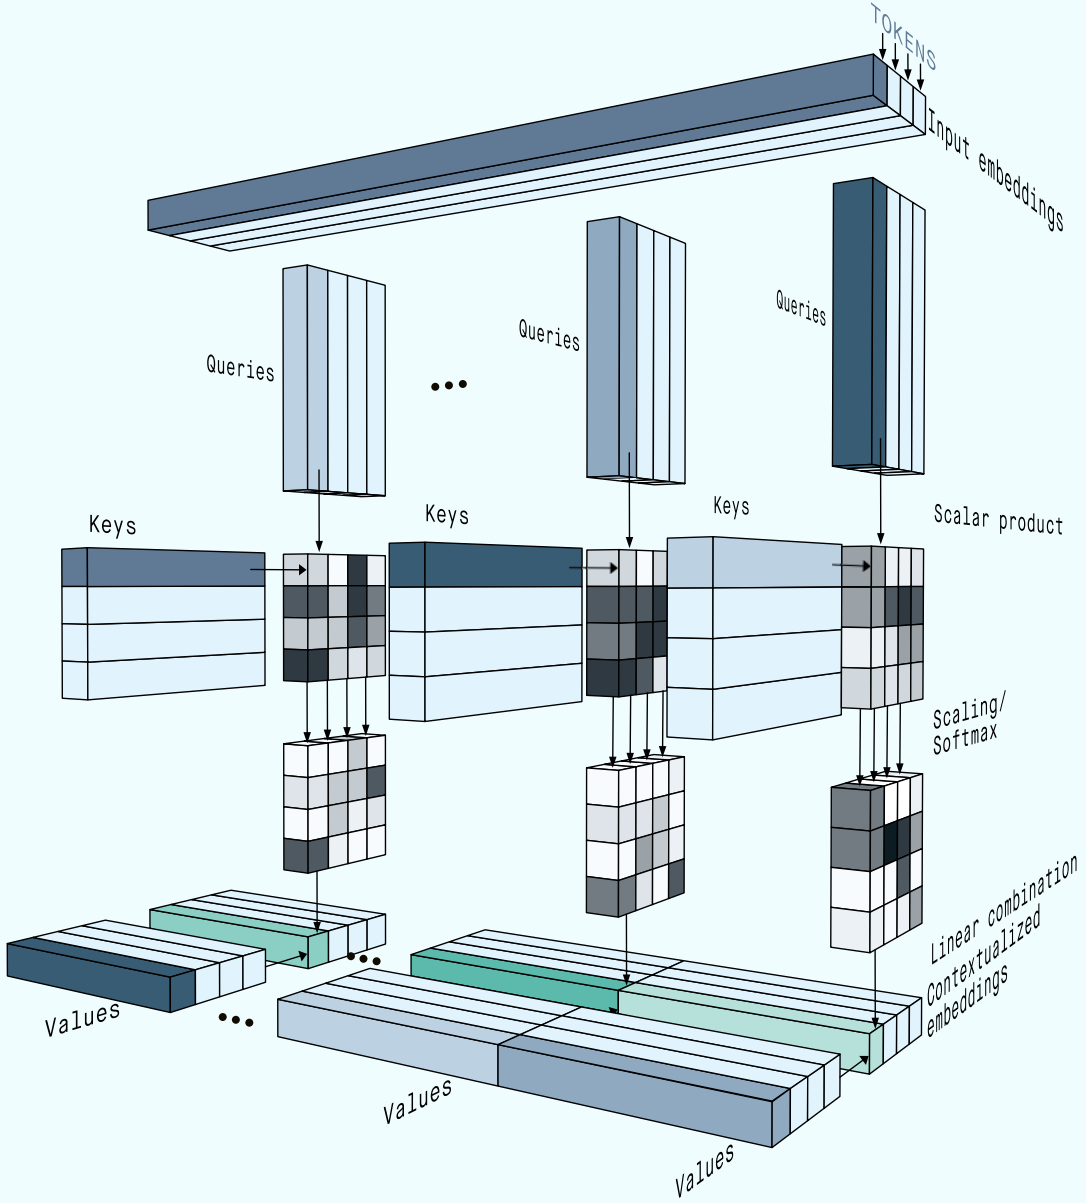

Multihead self-attention

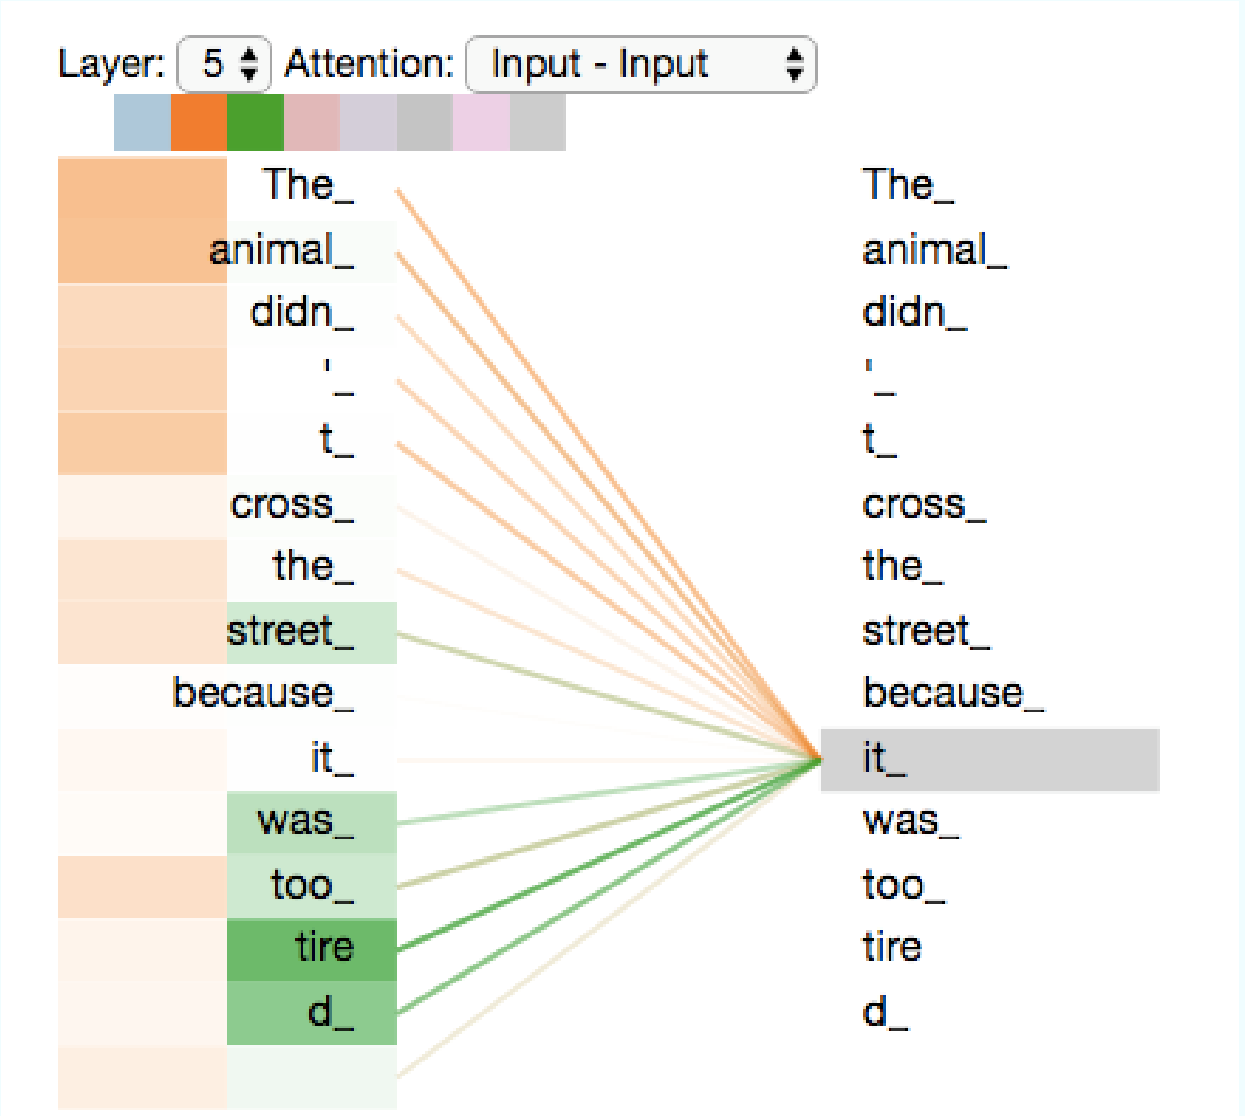

Positional encodings

$$\mathrm{PE}_{pos,2i} = \sin(pos/10000^{2i/d_{model}}),$$

$$\mathrm{PE}_{pos,2i+1} = \cos(pos/10000^{2i/d_{model}}),$$

Positional encodings

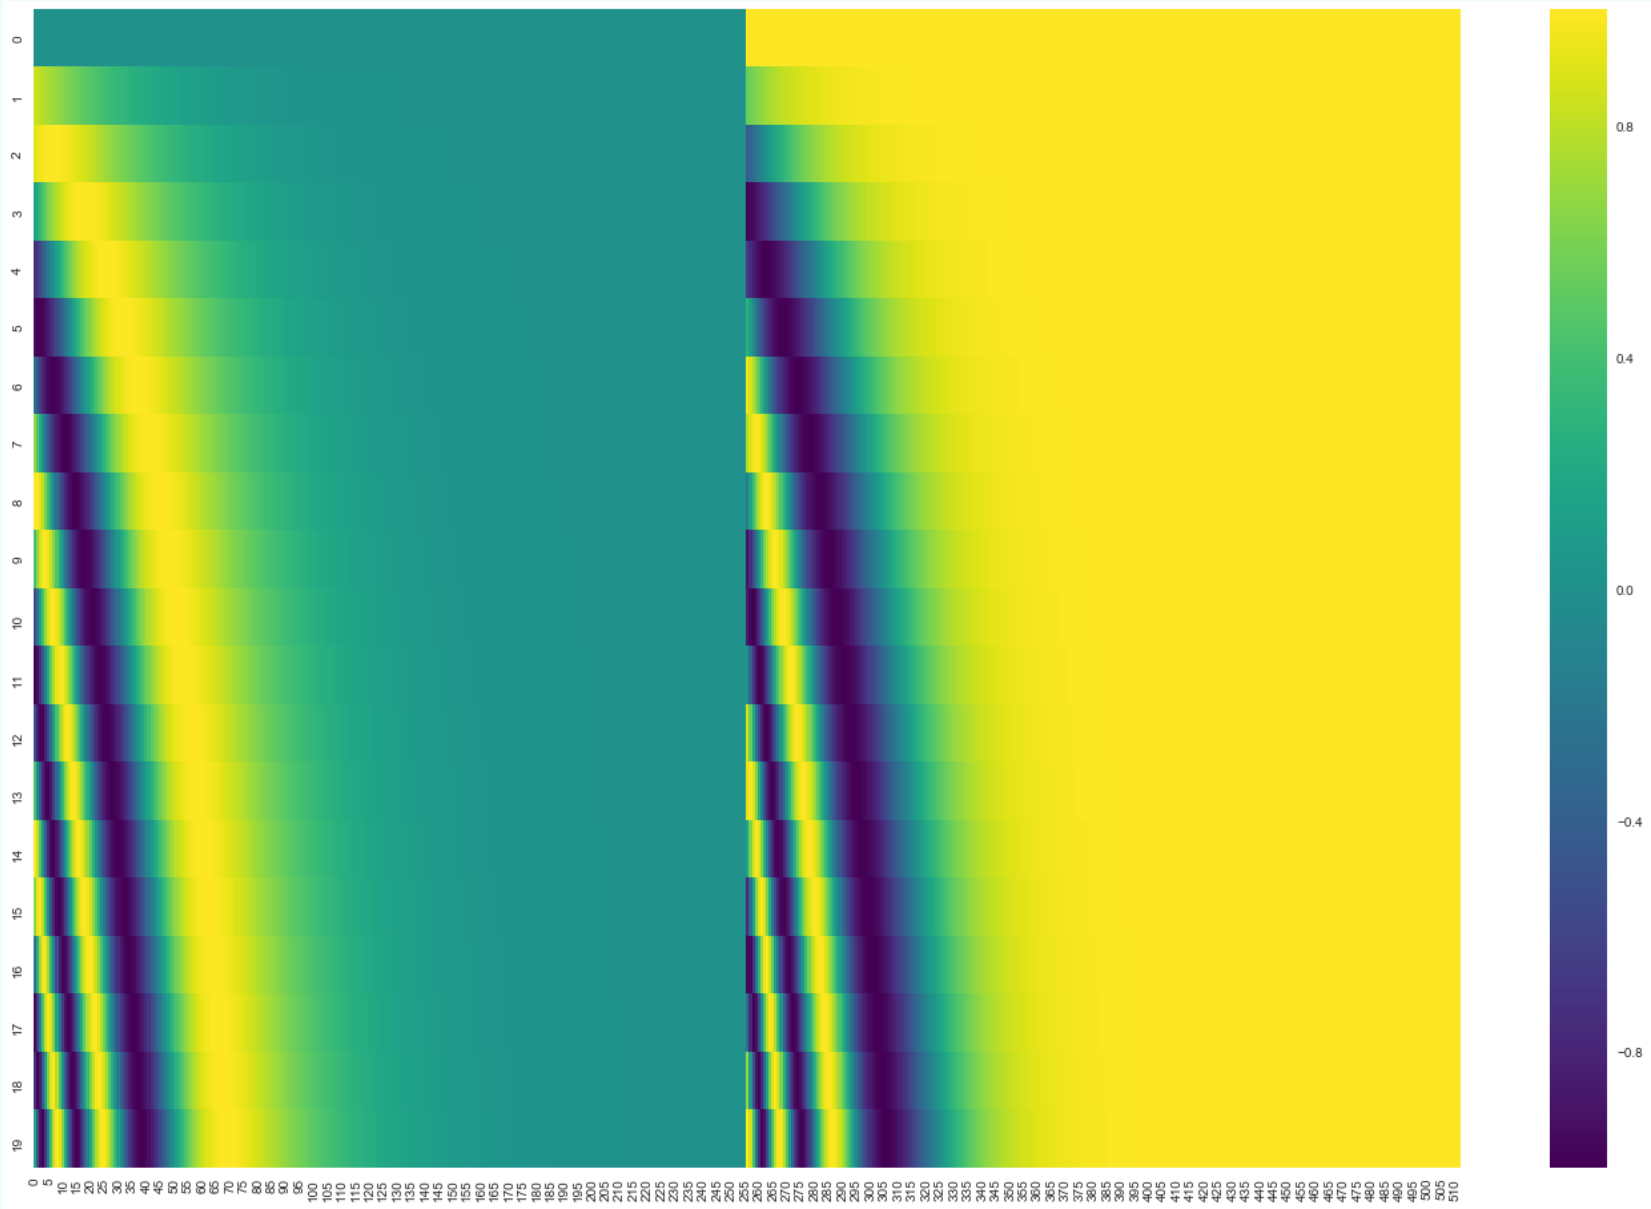

Transformer architecture

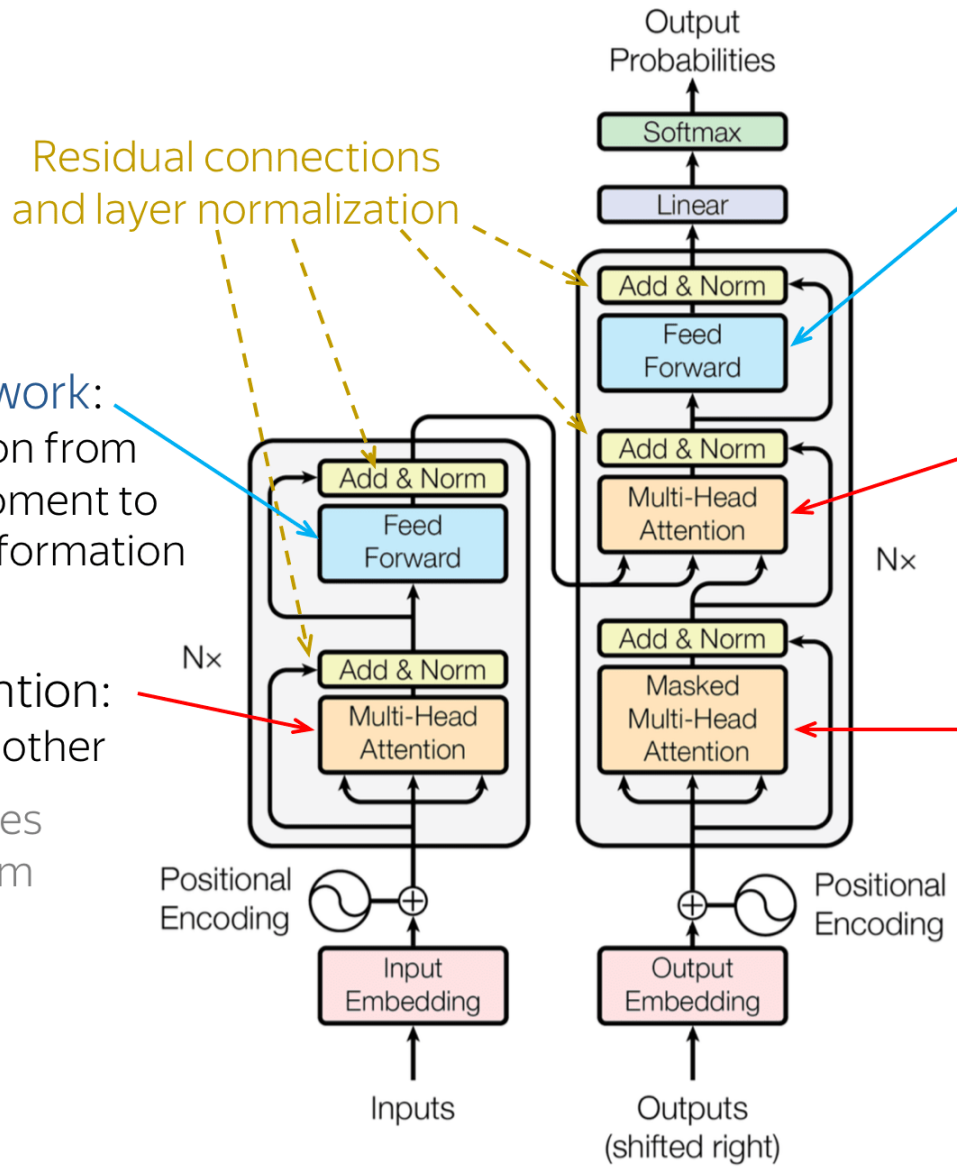

Residual connections and layer normalization

Feed-forward network: after taking information from other tokens, take a moment to think and process this information

Encoder self-attention: tokens look at each other

queries, keys, values are computed from encoder states

Feed-forward network: after taking information from other tokens, take a moment to think and process this information

Decoder-encoder attention: target token looks at the source

queries – from decoder states; keys and values from encoder states

Decoder self-attention (masked): tokens look at the previous tokens

queries, keys, values are computed from decoder states

AIRI

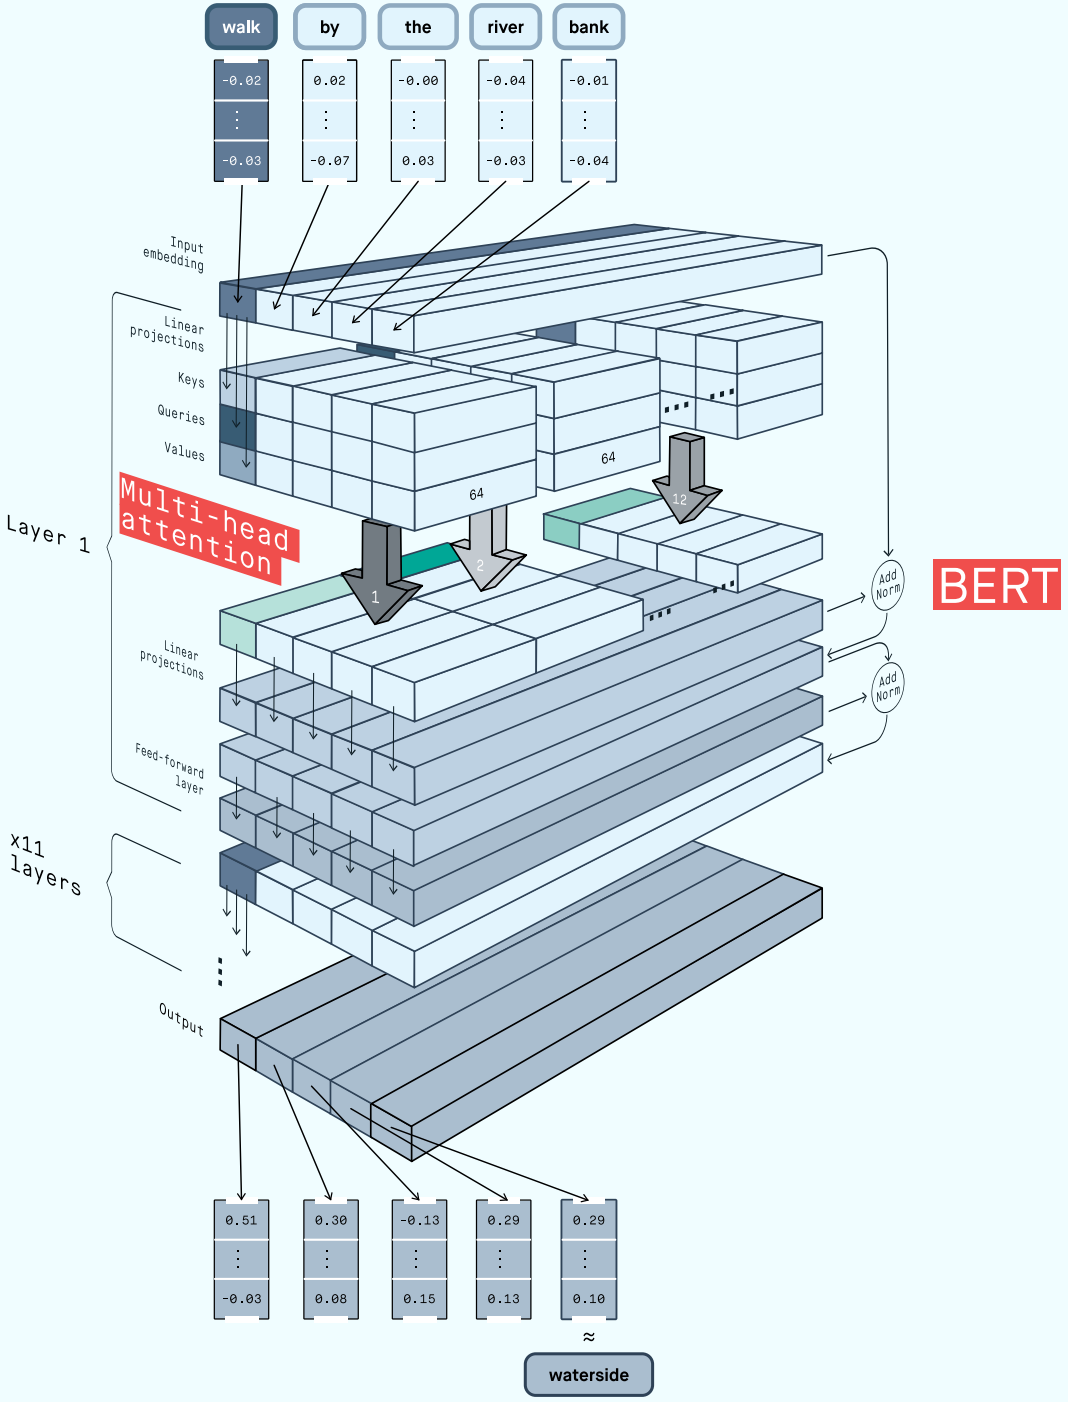

airi.net